

# **Introduction to MapReduce**

# MapReduce?

- A programming model or abstraction.
- It is not a programming language
- A novel way of thinking about designing a solution to certain big data problems...
- It enables us to
  1. Partition data into small chunks (called partitions)
  2. Execute tasks in parallel by:
    - Mappers
    - Reducers

# MapReduce as a Model/Paradigm/Architecture

Implementations of MapReduce:

- Google App Engine
- Apache Hadoop
- Apache Tez
- Apache Spark (implements superset of MapReduce)
- Snowflake
- Amazon Athena

# Motivation of MapReduce

- Process lots of data in parallel
  - Google processed about 24 petabytes of data per day in 2009.
  - Facebook processed 60 petabytes of data per day in 2020
- **A single machine** cannot serve all the data
  - You need a distributed system (called cluster computing) to store and process in parallel
- Parallel programming

# Motivation


- Parallel programming?
  - **Concurrency/Threading/Parallelism** is hard!
  - How do you facilitate **communication** between servers/nodes?
  - How do you **scale to more machines**?
  - How do you handle machine **failures**?
  - How do we handle disk **failures**?
- MapReduce uses cluster computing and enables parallelism

# What is a Cluster Computing?

- MapReduce uses cluster computing and enables parallelism
- A computer cluster is a set of computers (nodes or servers) that work together so that they can be viewed as a single system.
- Example: A cluster of 101 nodes:
  - One Master (as a cluster manager)
  - 100 worker nodes

# What is a Cluster Computing?

Perform tasks in parallel using a set of computers



# MapReduce reintroduced...

- Google created the awareness by publishing a paper
- *Apache Hadoop* made it into a sensation
- *Apache Hadoop* is an open-source MapReduce implementation based on Google's paper.
- *Spark implements superset of MapReduce*

MapReduce: Simplified Data Processing on Large Clusters

By Jeffrey Dean and Sanjay Ghemawat, Google

OSDI'04: Sixth Symposium on Operating System Design and Implementation.  
December, 2004.



# MapReduce reintroduced...

## Google's big problem:

- Index billions of web documents *everyday!*
- *Takes too much time and effort!*

## Solution:

*Use MapReduce to utilize 100's or 1000's of servers (cluster computing) to index billions of documents*

# (key, value) in MapReduce

For MapReduce, everything (input and output) is expressed as a tuple of 2 values: (key, value)

**NOTE: key and value can be any data types**

- **Mapper:**


- `map(key, value)`
  - key: as a partition number, or a record number
  - value: an actual input record
- Outputs { {k2, v2), ... }

- **Reducer:**

- `reduce(key2, value2)`
  - key2 : a k2 in mappers
  - value2 : Iterable<objects> OR list of objects
- Outputs: { (k3, v3), ... }

# Hadoop implements MapReduce

1. Partition input into small chunks
2. Parallelize chunks in many servers
3. Apply `map()` and `reduce()`



# Word Count: Notation

## **() denotes a Tuple**

- Example-1: `(1, 2, 4)` denotes tuple of 3 integers
- Example-2: `("fox", 3)` denotes a **(key, value)** pair  
Where **key** is "fox" and **value** is 3
- Example-3: `()` denotes an empty tuple

## **[] denotes a List of 0, 1, 2, ... items/objects**

- Example-1: `[1, 2, 6]` is a list of 3 integers
- Example-2: `["to", "be"]` is a list of 2 strings
- Example-3: `[]` denotes an empty list

## **{ } denotes a Set of 0, 1, 2, ... or more objects**

- Example-1: `{ ("a", 2), ("b", 3), ("z", 9) }`  
denotes a set of 3 pair objects
- Example-2: `{ }` denotes an empty set

# Word Count: Notation

**Iterable<object>** denotes a list of objects

- Example-1: **Iterable<Integer>**
  - [1, 2, 2, 3, 4]
  - [1, 1, 1, 1, 1, 1, 1]
- Example-2: **Iterable<String>**
  - ["a", "fox", "jumped"]
  - ["a", "a", "b", "b", "b"]
- Example-3: **Iterable<(String, Integer)>**
  - [("a", 10), ("fox", 7), ("jumped", 8)]
  - [("key1", 10), ("key200", 700)]

# MapReduce flow... 1

1. Input is partitioned and passed to mappers
2. Mappers get input as (**key, value**) pairs
  - **key** might be a partition number or record number and might be ignored (if not needed)
  - **value:** as an input record (as a String)
  - `map(key, value):` emits a set of `{(key2, value2)}` pairs
3. Output of mappers is passed into **Sort & Shuffle** system (provided by MapReduce implementation)

## MapReduce flow... 2

4. Assume that output from all mappers are:

(there are N unique keys: {Key\_1, Key\_2, ..., Key\_N})

(Key\_1, v\_11), (Key\_1, v\_12), ...

(Key\_2, v\_21), (Key\_2, v\_22), ...

...

(Key\_N, v\_N1), (Key\_N, v\_N2), ...

Then Sort & Shuffle groups values of mappers by their associated keys.  
Sort & Shuffle outputs (key, value) pairs as:

(Key\_1, [v\_11, v\_12, ...])

(Key\_2, [v\_21, v\_22, ...])

...

(Key\_N, [v\_N1, v\_N2, ...])

Sort & Shuffle  $\leftrightarrow$  SQL's GROUP BY

## MapReduce flow... 3

5. Output of Sort & Shuffle is passed to reducers:

```
(Key_1, [v_11, v_12, ...])
```

```
(Key_2, [v_21, v_22, ...])
```

...

6. A reducer will operate/execute on (key, value) produced by Sort & Shuffle

```
# key : one of Key_1 or Key_2, ...
```

```
# values: associated values for a given key Key_1 or Key_2, ...
```

```
reduce(key, values){
```

```
    <reducer's logic, which may emit any number  
    of (K3, V3) pairs>
```

```
}
```



# MapReduce Example: word count

MapReduce works with (key, value) pairs

- Mappers Input: as (key, value) pair
- `map(123, "fox jumped and fox jumped and jumped and jumped")`
  - key: 123 as a key is a record number (ignored here)
  - value: "fox jumped and fox jumped and jumped and jumped"

Mappers output as (K, V) pairs: K is a word, V is a frequency

(fox, 1), (jumped, 1), (and, 1),  
(fox, 1), (jumped, 1), (and, 1),  
(jumped, 1), (and, 1), (jumped, 1)

## MapReduce Example continued...

### Mappers output:

```
(fox, 1), (jumped, 1), (and, 1),  
(fox, 1), (jumped, 1), (and, 1),  
(jumped, 1), (and, 1), (jumped, 1)
```

### Sort & Shuffle output:

GROUP BY mapper's output KEY

```
(fox, [1, 1])  
(jumped, [1, 1, 1, 1])  
(and, [1, 1, 1])
```

## MapReduce Example continued...

Sort&Shuffle output: (used as input to reducers)

(fox, [1, 1])

(jumped, [1, 1, 1, 1])

(and, [1, 1, 1])

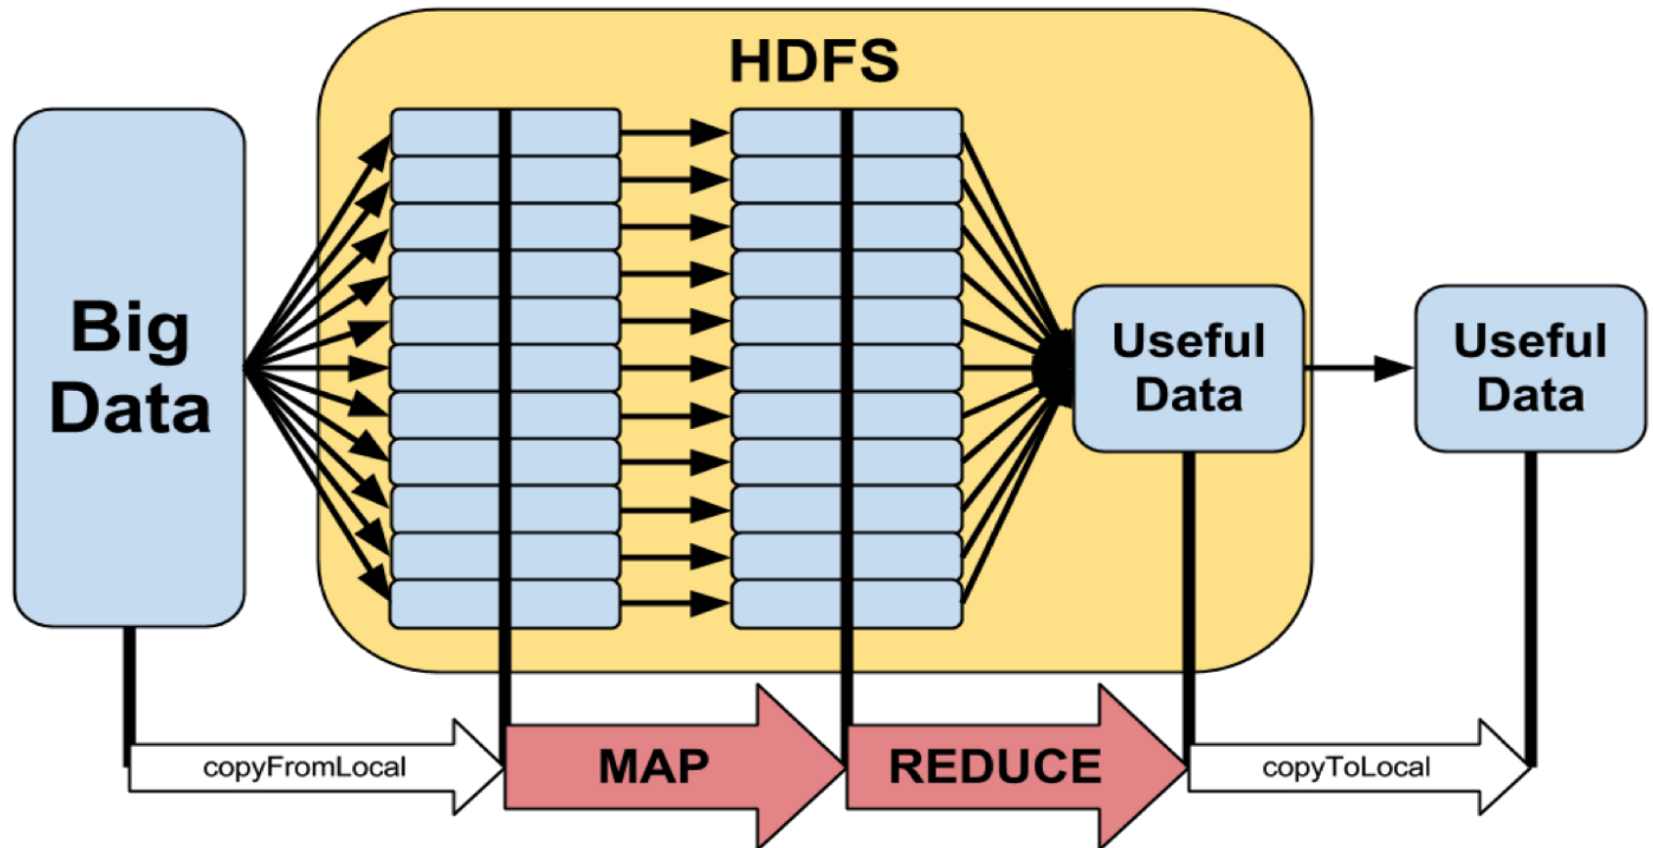
Reducers output:

(fox, 2)

(jumped, 4)

(and, 3)

# MapReduce Model: Hadoop Implementation



## MapReduce provides:

- Partition data into small chunks
- Automatic parallelization, distribution
- I/O scheduling
- Load balancing
- Network and data transfer optimization
- Fault tolerance
  - Handling of machine failures

## MapReduce provides:

- Fault tolerance
  - Handling of machine failures
  - Fault tolerance refers to the ability of a system (computer, network, cloud cluster, etc.) to continue operating without interruption when one or more of its components (disk, computer, ...) fail.

# MapReduce: Scale-Out, but do NOT Scale-Up

- Need more power: Scale-Out
  - Large number of commodity servers
  - Not expensive to add or replace servers
  - Add more servers any time
- Do NOT Scale-Up
  - Do NOT use high end specialized servers
  - Very Expensive
  - Replacement very costly

# MapReduce: Scale-Out, but do NOT Scale-Up

- Need more power: Scale-Out
  - Large number of commodity servers
  - Not expensive to add or replace servers
  - Add more servers any time
- A commodity server is a commodity computer that is dedicated to running server programs and carrying out associated tasks. In many environments, multiple low-end servers share the workload. Commodity servers are often considered disposable and, as such, are replaced rather than repaired.



# MapReduce Implementations

## **Google App Engine**

(proprietary, not open-source)

## **Apache Hadoop: implementation of MapReduce**

(open-source, based on Google's paper)

## **Apache Spark:**

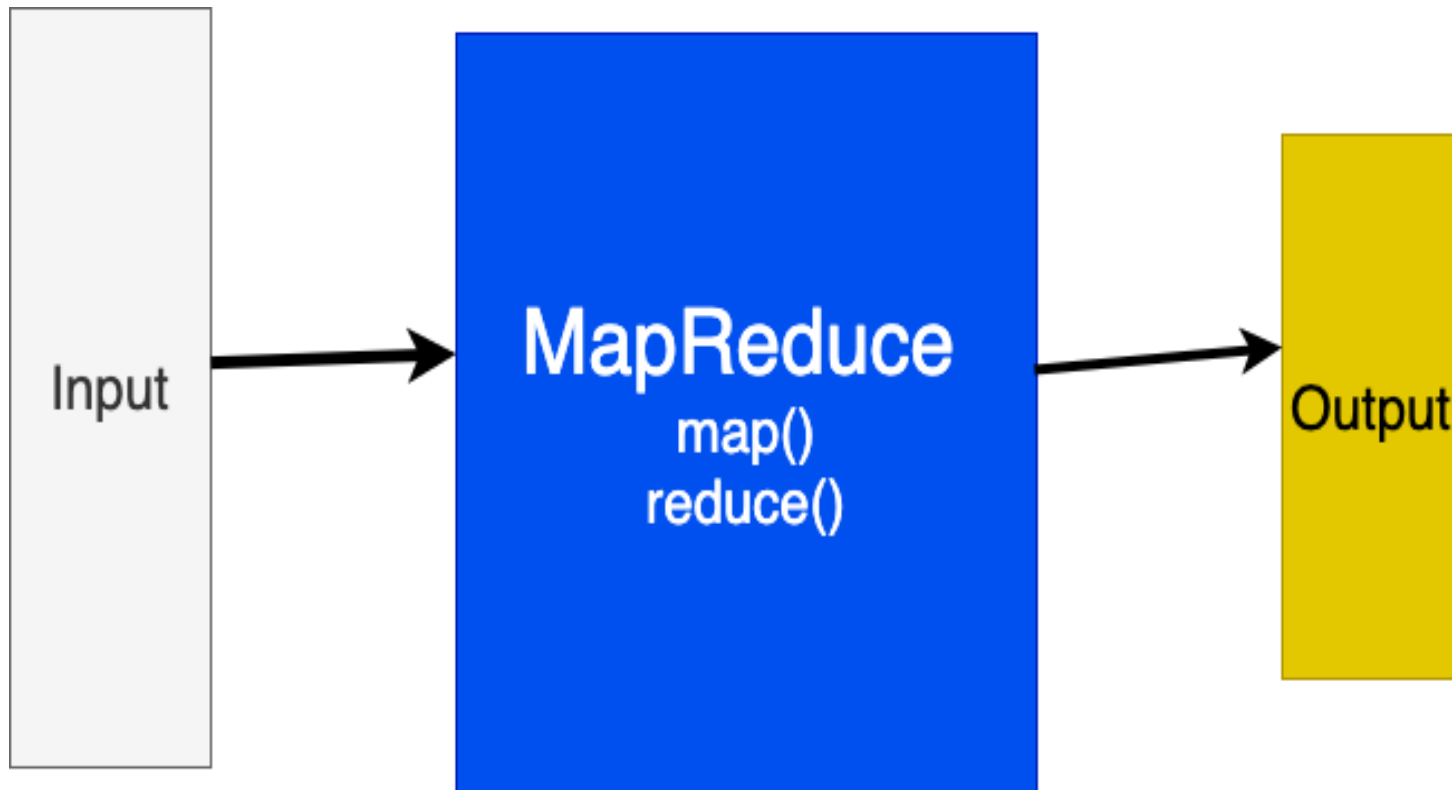
open-source, superset of MapReduce

In-memory computing, very fast

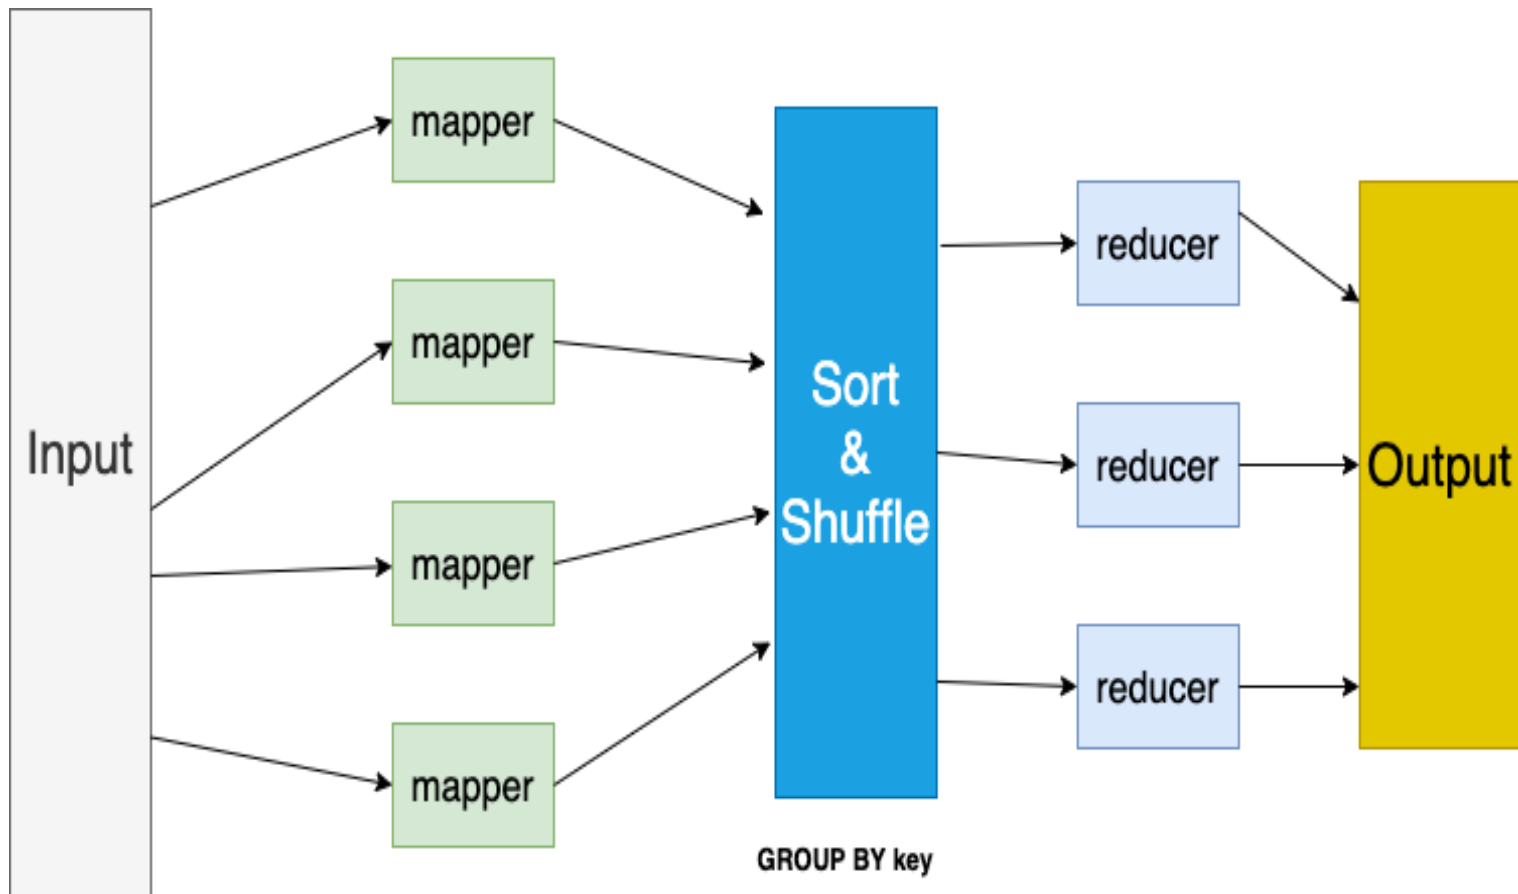
# Typical problem solved by MapReduce

- Read a lot of data
- Partition data into small chunks
- **map(key, value):** extract something you care about from each record
  - output: set of (key2, value2) pairs
- **Shuffle and Sort** [done by MapReduce Implementation]
  - Output as (key2, [value\_21, value\_22, value\_23, ...])
- **reduce(key2, values):** aggregate, summarize, filter, or transform
  - Output: set of (key3, value3) pairs
- Write the results

# MapReduce model



# MapReduce model



# Mappers in parallel

- **map() function** run in parallel.
- Each mapper operates on a set of chunks assigned to it by the job tracker.
- Mappers write to local disk.

# key: a record number or hash of a record

# value: a single record (actual data)

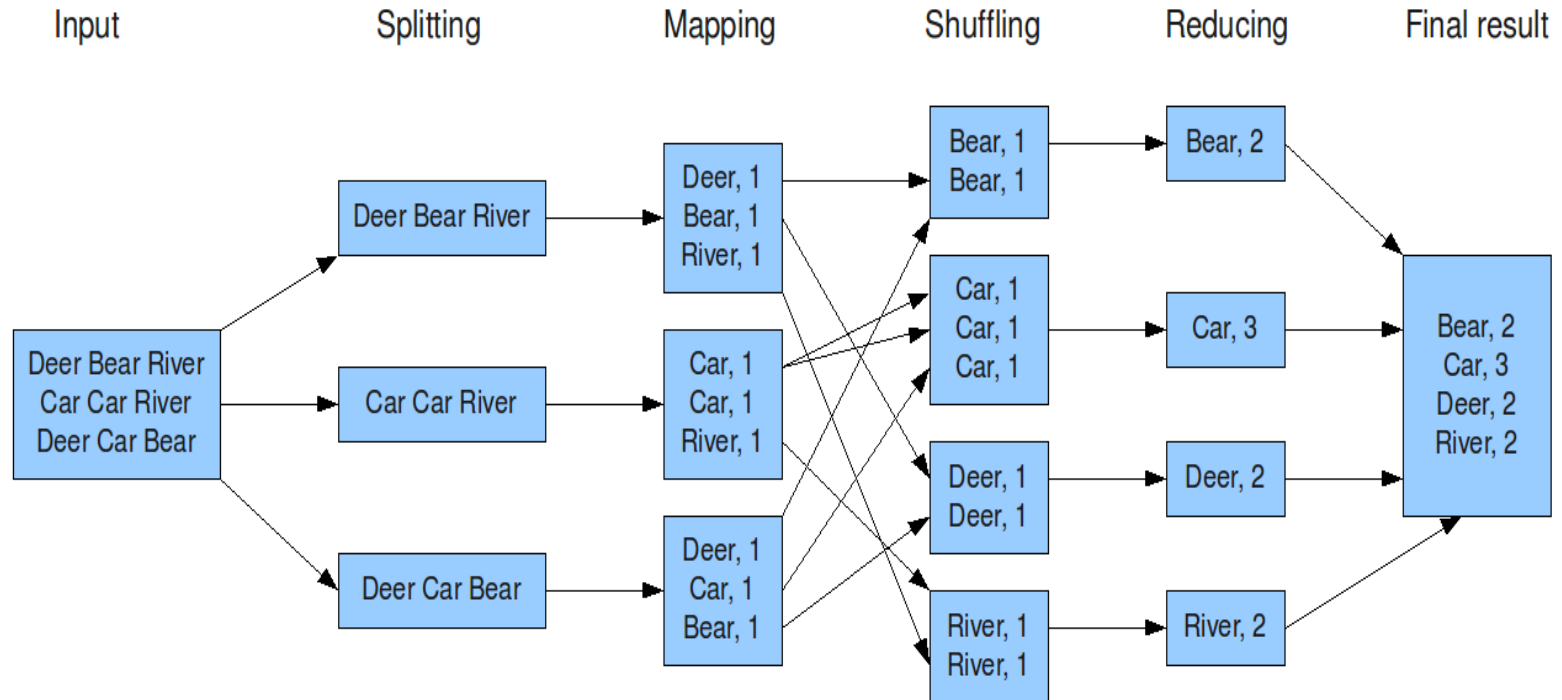
**map(key, value):**

can emit any number of (K, V) pairs:

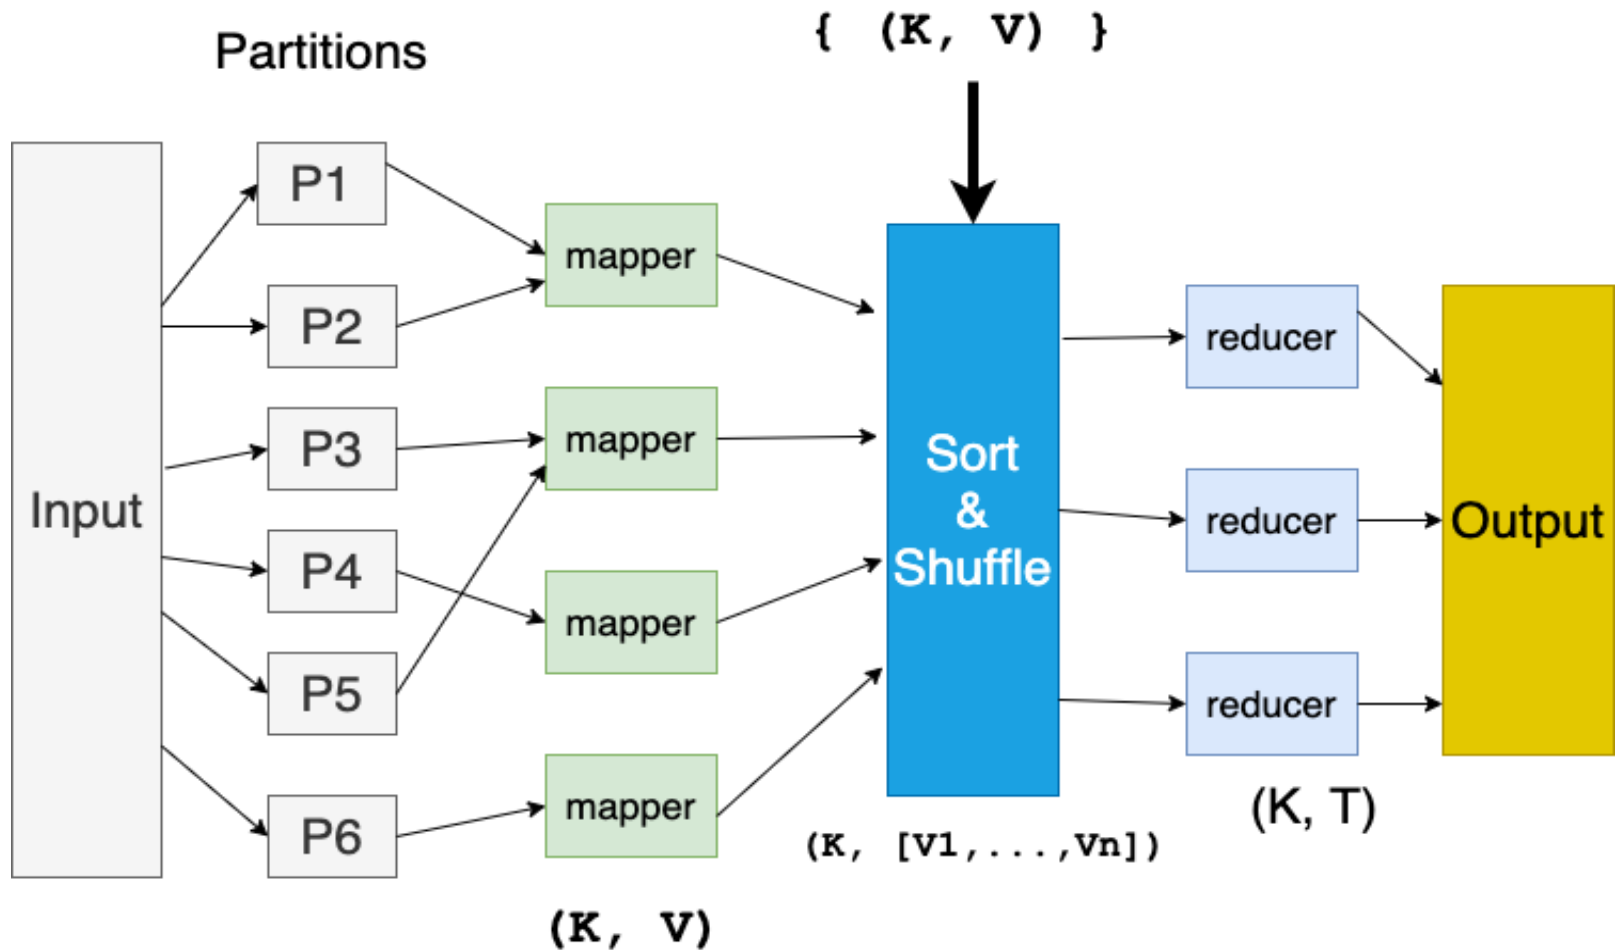
(K\_1, V\_1), (K\_2, V\_2), ... (K\_n, V\_n)

# MapReduce model

The overall MapReduce word count process



# MapReduce model



# Reducers in parallel

- **reduce() function** run in parallel.
- Each reducer() operates on  
(key, [V\_1, V\_2, ..., V\_n])

- reduce((key, [V\_1, V\_2, ..., V\_n]):

Reducers create final outputs as:


- (K1, T1)
- (K2, T2)
- ...



# MapReduce workflow

Input Data

Output Data



## Map

extract something you  
care about from each  
record

## Reduce

aggregate,  
summarize, filter, or  
transform

# Mappers and Reducers

- Need to handle **more data**?
  - Just add **more Mappers/Reducers**!
- No need to handle **multithreaded code** 😊
  - Mappers and Reducers are typically single threaded and **deterministic**
    - **Determinism** allows for **restarting of failed jobs**
  - Mappers/Reducers run **entirely independent** of each other
  - Mappers/Reducers run in **separate JVM processes**

## Word Count Problem: solve it by MapReduce

- INPUT: Given a set of text documents/files
- OUTPUT: Find frequencies of each unique word
- Input: “gray fox red fox jumped over red gray fox”
- Output: dictionary[(word, frequency)]
  - (gray, 2)
  - (red, 2)
  - (fox, 3)
  - (jumped, 1)
  - (over, 1)

# Mapper: map (key, value)

- Reads in input pair as (Key, Value)
- Outputs a set of pairs (K', V')
  - Let's count number of each word in user queries (or Tweets/Blogs)
  - The input to the mapper will be (Query\_ID, Query\_Text):  

```
<Q1, "the teacher went to the store. the store was  
closed; the store opens in the morning. the store  
opens at 9am." >
```
  - The mappers output would be:  

```
(the, 1) (teacher, 1) (went, 1) (to, 1) (the, 1)  
(store, 1) (the, 1) (store, 1) (was, 1) (closed, 1)  
(the, 1) (store, 1) (opens, 1) (in, 1) (the, 1)  
(morning, 1) (the, 1) (store, 1) (opens, 1) (at, 1)  
(9am, 1)
```

## Sort & Shuffle: SQL's **GROUP BY**

- Accepts the **Mapper output**, and aggregates values on the key

- For our example, the mappers output would be:

(the, 1) (teacher, 1) (went, 1) (to, 1) (the, 1) (store, 1) (the, 1)  
(store, 1) (was, 1) (closed, 1) (the, 1) (store, 1) (opens, 1) (in, 1)  
(the, 1) (morning, 1) (the, 1) (store, 1) (opens, 1) (at, 1) (9am, 1)

- The output of Sort & Shuffle would be:

(the, [1, 1, 1, 1, 1, 1])

(store, [1, 1, 1])

...

**Reducer:** `reduce(key, values)`

- Accepts the Sort & Shuffle output,
- and aggregates values on the key
  - Input to reducer: (`store`, [1, 1, 1])
    - Output: (`store`, 3)
  - Input to reducer : (`the`, [1, 1, 1, 1, 1, 1])
    - Output: (`the`, 6)
  - ...

# MapReduce Job Components

1. Input path (identify your input files)
2. Output path (where to write output)
3. `map()` function

**`map(key, value)`**

`emits { (K2, V2), ... }`

- 3.5 [Sort & Shuffle is done by MapReduce]

4. `reduce()` function

**`reduce(K2, [value_1, value_2, ...])`**

`emits { (K3, V3), ... }`

5. OPTIONAL `combine()` function

## MapReduce Job: **Input Path**

- **Input path:** 3 files will be read
- **Example:** `s3://my_bucket/project7/`

`s3://my_bucket/project7/file1.txt`

`s3://my_bucket/project7/file2.txt`

`s3://my_bucket/project7/file3.txt`



## MapReduce Job: **Output Path**

**Reducers output will be written to output path**

- **Output path:**

- **Example:** `s3://my_bucket/output7/`

`s3://my_bucket/output7/_SUCCESS`

`s3://my_bucket/output7/part1`

`s3://my_bucket/output7/part2`

`s3://my_bucket/output7/part3`

`s3://my_bucket/output7/part4`

# MapReduce Job: map() function for Word Count

# **pseudo-code**

# key: may be a record number and ignored here

# value: a single record of your input data

# "fox jumped and jumped"

map(key, value) {

# tokenize record

words = value.split(" ")

# array index: 0            1            2            3

# words = ["fox", "jumped", "and", "jumped"]

for word in words {

emit (word, 1)

}

}

# MapReduce Job: map() function

```
# pseudo-code
# key: may be a record number and ignored here
# value: a single record of your input data
# "fox jumped and jumped"
map(key, value) {
    # tokenize record
    words = value.split(" ")

    # array index: 0      1      2      3
    # words = ["fox", "jumped", "and", "jumped"]
    for word in words {
        emit (word, 1)
    }
}
```

## Output of a mapper:

```
(fox, 1)
(jumped, 1)
(and, 1)
(jumped, 1)
```

## MapReduce Job: map() function

# 103 is a record number

INPUT: (103, “a fox of jumped over red fox and jumped”)

### Mappers output:

(a, 1)

(fox, 1)

(of, 1)

(jumped, 1)

(over, 1)

(red, 1)

(fox, 1)

(and, 1)

(jumped, 1)

**MapReduce Job: map() function**

**FILTER: Ignore words with length of less than 3 Chars.**

```
# pseudo-code
# key: may be a record number
# value: the entire record such
# as "fox jumped and jumped"
map(key, value) {
    words = value.split(" ")
    for word in words {
        # filter non-desired words
        if (len(word) > 2) {
            emit( word, 1)
        }
    }
}
```

**MapReduce Job: map() function**

**Ignore words with length of less than 3 Chars.**

INPUT: (103, “a fox of jumped over red fox and jumped”)

**Mappers output:**

(fox, 1)

(jumped, 1)

(over, 1)

(red, 1)

(fox, 1)

(and, 1)

(jumped, 1)

NOTE: FILTERING: “a” and “of” were dropped  
since their length is less than 3

## MapReduce Job: map() function

**FILTER-1:** Ignore records with length of less than 80 chars

**FILTER-2:** ignore words less than 3 characters

# **pseudo-code**

# key: may be a record number

# value: the entire record such as "fox jumped and jumped"

```
map(key, value) {  
    # filter-1  
    if (len(value) < 80) {  
        # no (K, V) is emitted at all  
        return  
    }  
    words = value.split(" ")  
    for word in words {  
        # filter-2  
        if (len(word) > 2) {  
            emit( word, 1) # filter non-desired words  
        }  
    }  
}
```

# MapReduce Job: map() function

# key: may be a record number such as 100 (is ignored)

# value: the entire record such as

# (103, “fox jumped and jumped”)

map (key, value) output:

(fox, 1)

(jumped, 1)

(and, 1)

(jumped, 1)



# MapReduce Job: map() function

# key: may be a record number such as 1234 (is ignored)

# value: the entire record such as

# (1234, "fox jumped over fox")

map(key, value) output:

(fox, 1)

(jumped, 1)

(over, 1)

(fox, 1)

# Sort & Shuffle: Receives Mappers Output

- Sort & Shuffle is the genie of MapReduce
- Similar to SQL's "GROUP BY"
- Output of Sort & Shuffle:

(key\_1, [V\_11, V\_12, ...])

(key\_2, [V\_21, V\_22, ...])

...

(key\_N, [V\_N1, V\_N2, ...])

- Therefore, N reducers are required.

# Input to Reducers

```
(key_1, [V_11, V_12, ...])
```

```
(key_2, [V_21, V_22, ...])
```

```
...
```

```
(key_N, [V_N1, V_N2, ...])
```

NOTE:

- All keys {key\_1, key\_2, ..., key\_N} are unique
- A key may have any number of values
- Values are not sorted (they can be in any order)

# MapReduce Job Components: reduce()

- Reducers receive output of Sort & Shuffle phase.
- `reduce(key, values)` accepts a
  - `key`: a single unique key
  - `values`: `[V_1, V_2, ..., V_n]`
- Creates any number of new  $(K', V')$  pairs

# MapReduce Job: reduce() function for Word Count

```
# pseudo-code
# handle (key, [v1, v2, ...])
# key: is a unique word
# values: Iterable<Integer>
#       such as [1, 1, ..., 1]
reduce(key, values) {
    count = 0
    # iterate values
    for v in values {
        count += v
    }
    emit(key, count)
}
```

MapReduce Job: reduce() function +  
FILTER: ignore words with frequencies of less than 5

```
# pseudo-code
# key: is a unique word
# values: [1, 1, ..., 1]
reduce(key, values) {
    count = 0
    for v in values {
        count += v
    }
    # filter
    if (count >= 5) {
        emit(key, count)
    }
}
```

**MapReduce Job: reduce() function + what if we want to ignore words with frequencies of less than 5**  
**Ignore words with length of less than 3 Chars.**

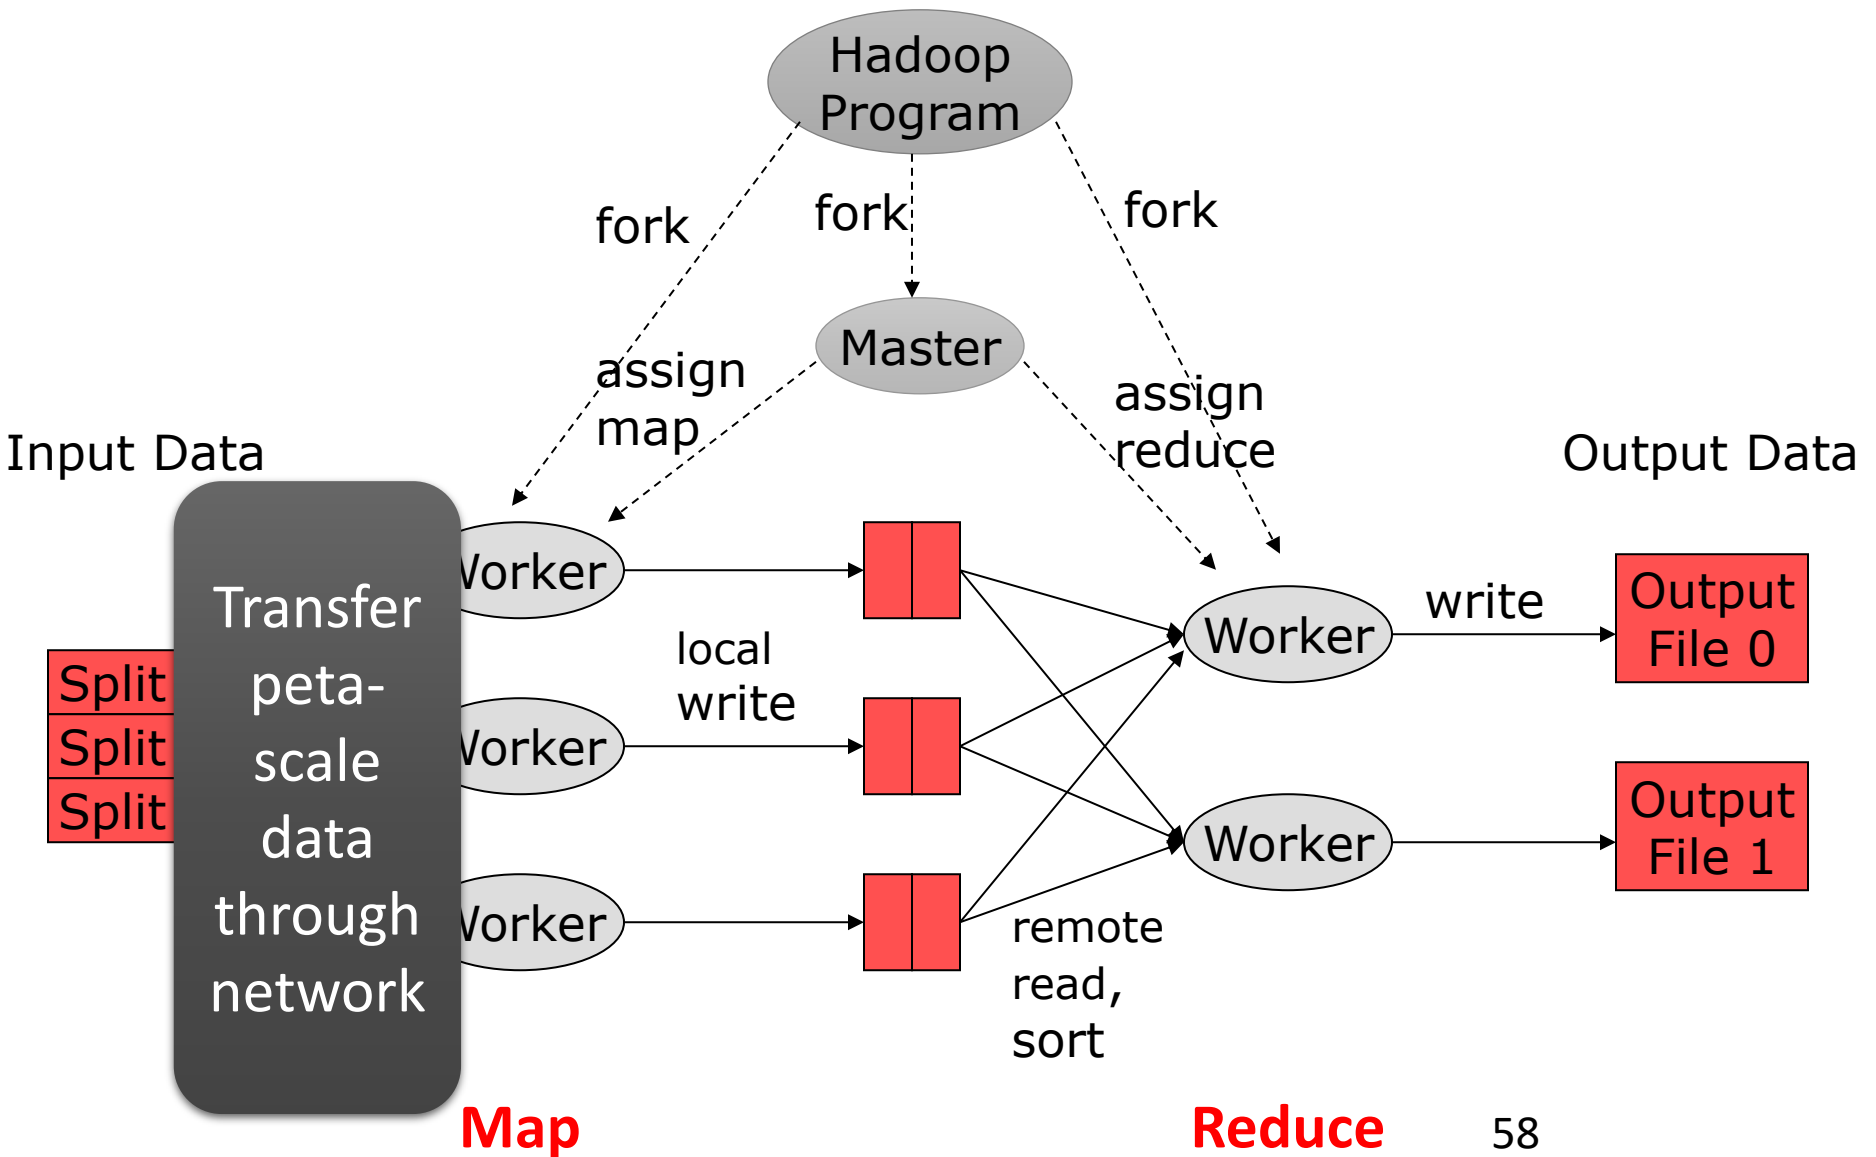
```
# pseudo-code
# key: is a unique word
# values: [1, 1, ..., 1]
reduce(key, values) {
    # not a proper filter for reducer (should be done in mapper)
    if (len(key) <= 2) {
        return
    }
    count = 0
    for v in values {
        count += v
    }
    # proper filter
    if (count < 5) {
        return
    }
    else {
        emit(key, count)
    }
}
```

## MapReduce Job: reduce() function and output length of key

```
# pseudo-code
# key: is a unique word
# values: [1, 1, ..., 1]
reduce(key, values) {
    count = 0
    for v in values {
        count += v
    }
    # create a new composite value
    new_value = (len(key), count)
    emit(key, new_value)
}
```



# MapReduce



# Failure in MapReduce

- **Failures** are **norm** in commodity hardware
- **Worker failure**
  - Detect failure via periodic **heartbeats**
  - **Re-execute** in-progress map/reduce tasks
- **Master failure**
  - Single point of failure; Resume from Execution Log
- **Robust**
  - Google's experience: **lost 1600 of 1800 machines once!**, but **finished fine**

```
public class WordCount {
```

```
    public static class Map extends MapReduceBase implements  
        Mapper<LongWritable, Text, Text, IntWritable> {  
        private final static IntWritable one = new IntWritable(1);  
        private Text word = new Text();
```

```
        public void map(LongWritable key, Text value, OutputCollector<Text, IntWritable>  
            output, Reporter reporter) throws IOException {  
            String line = value.toString();  
            StringTokenizer tokenizer = new StringTokenizer(line);  
            while (tokenizer.hasMoreTokens()) {  
                word.set(tokenizer.nextToken());  
                output.collect(word, one);  
            }  
        }  
    }  
}
```

**Mapper**

```
    public static class Reduce extends MapReduceBase implements
```

```
        Reducer<Text, IntWritable, Text, IntWritable> {  
        public void reduce(Text key, Iterator<IntWritable> values, OutputCollector<Text,  
            IntWritable> output, Reporter reporter) throws IOException {  
            int sum = 0;  
            while (values.hasNext()) { sum += values.next().get(); }  
            output.collect(key, new IntWritable(sum));  
        }  
    }  
}
```

**Reducer**

```
public static void main(String[] args) throws Exception {  
    JobConf conf = new JobConf(WordCount.class);  
    conf.setJobName("wordcount");  
    conf.setOutputKeyClass(Text.class);  
    conf.setOutputValueClass(IntWritable.class);  
    conf.setMapperClass(Map.class);  
    conf.setCombinerClass(Reduce.class);  
    conf.setReducerClass(Reduce.class);  
    conf.setInputFormat(TextInputFormat.class);  
    conf.setOutputFormat(TextOutputFormat.class);  
    FileInputFormat.setInputPaths(conf, new Path(args[0]));  
    FileOutputFormat.setOutputPath(conf, new Path(args[1]));  
  
    JobClient.runJob(conf);  
}
```

Run this program as  
a MapReduce job

# Complete Word Count Solution in MapReduce/Hadoop

## Word Count Solution in MapReduce/Hadoop

- Driver Program
- Mapper Program
- Reducer Program

## Summary: MapReduce

- Programming paradigm for data-intensive computing
- Distributed & parallel execution model
- Simple to program
- The MapReduce framework automates many tedious tasks:
  - Data partitioning
  - Machine selection,
  - Failure handling
  - Sort & Shuffle

# Word Count in Hadoop/MapReduce

- We will NOT study Hadoop
- For reference, I am including a pointer to MapReduce Tutorial in Hadoop, which can help you to understand an implementation of MapReduce.